

# Recap & Look ahead

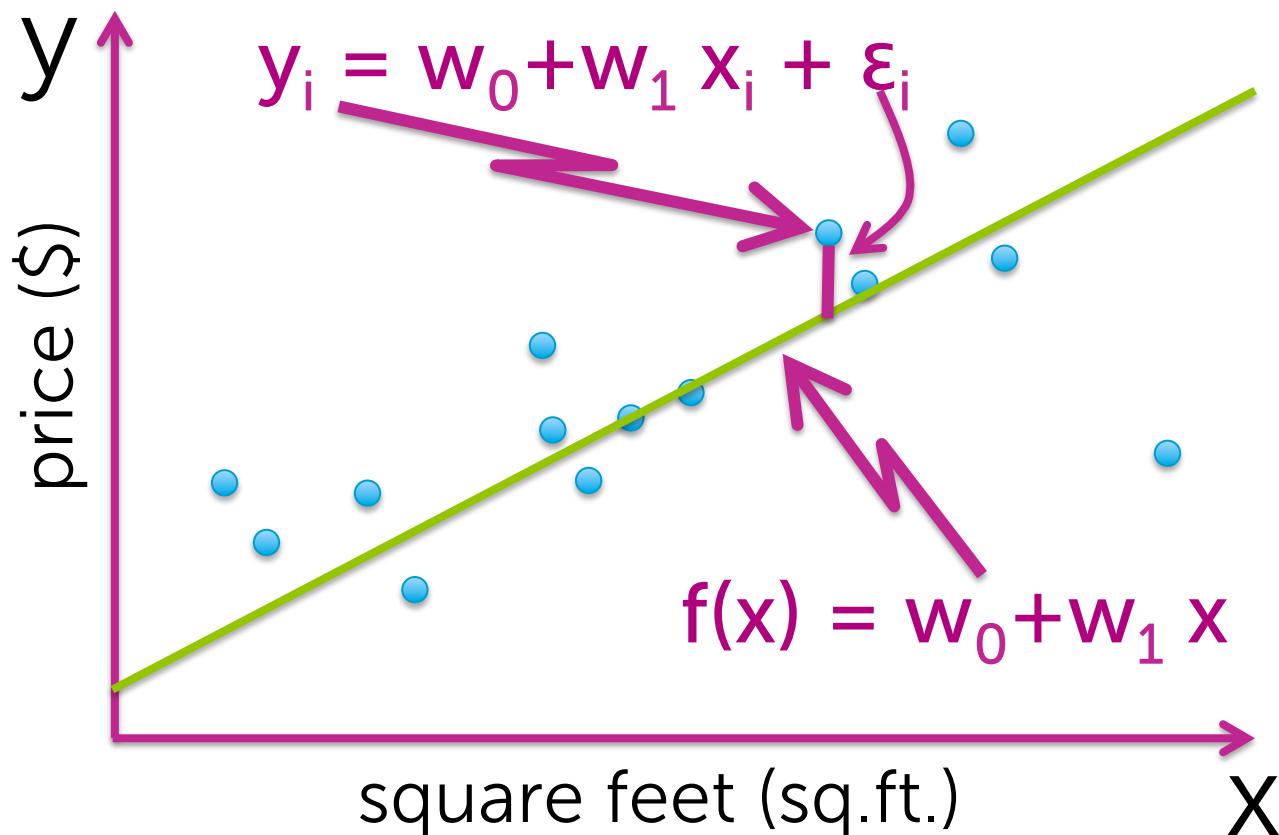
Emily Fox & Carlos Guestrin  
Machine Learning Specialization  
University of Washington

# What we've learned

# Module 1: Simple Regression

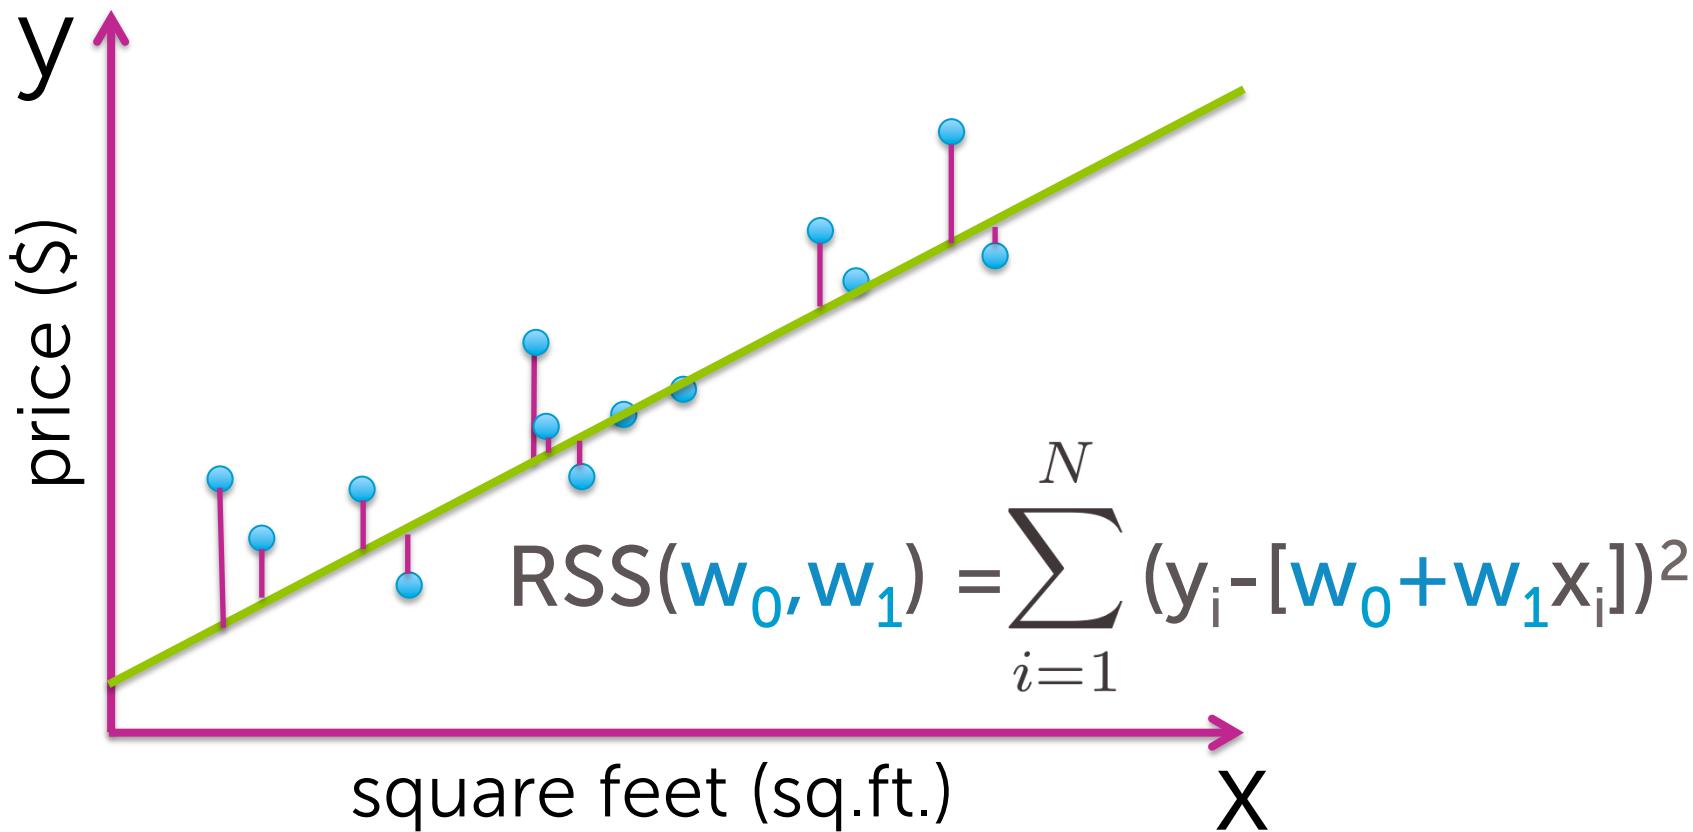
# Simple linear regression model

1 input and just fit a line to data



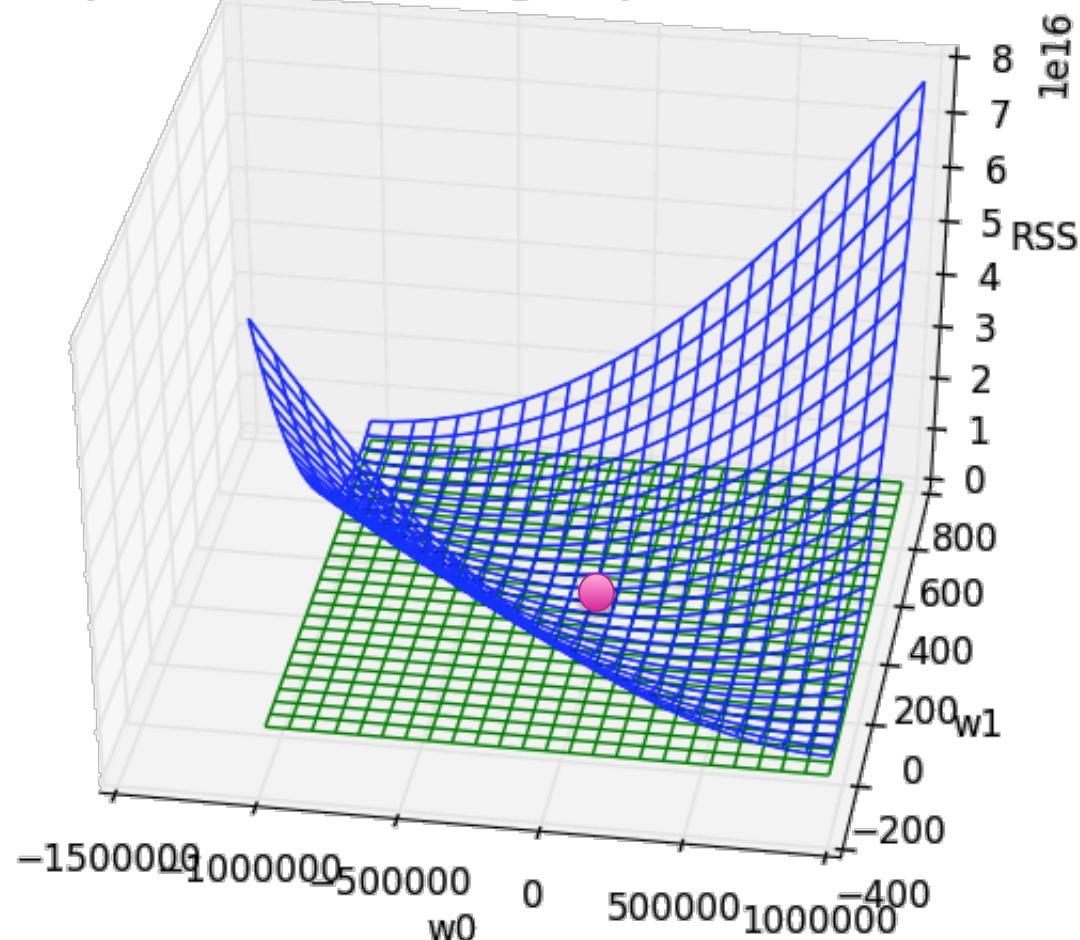
# "Cost" of using a given line

Residual sum of squares (RSS)



# Minimizing the cost

3D plot of RSS with tangent plane at minimum

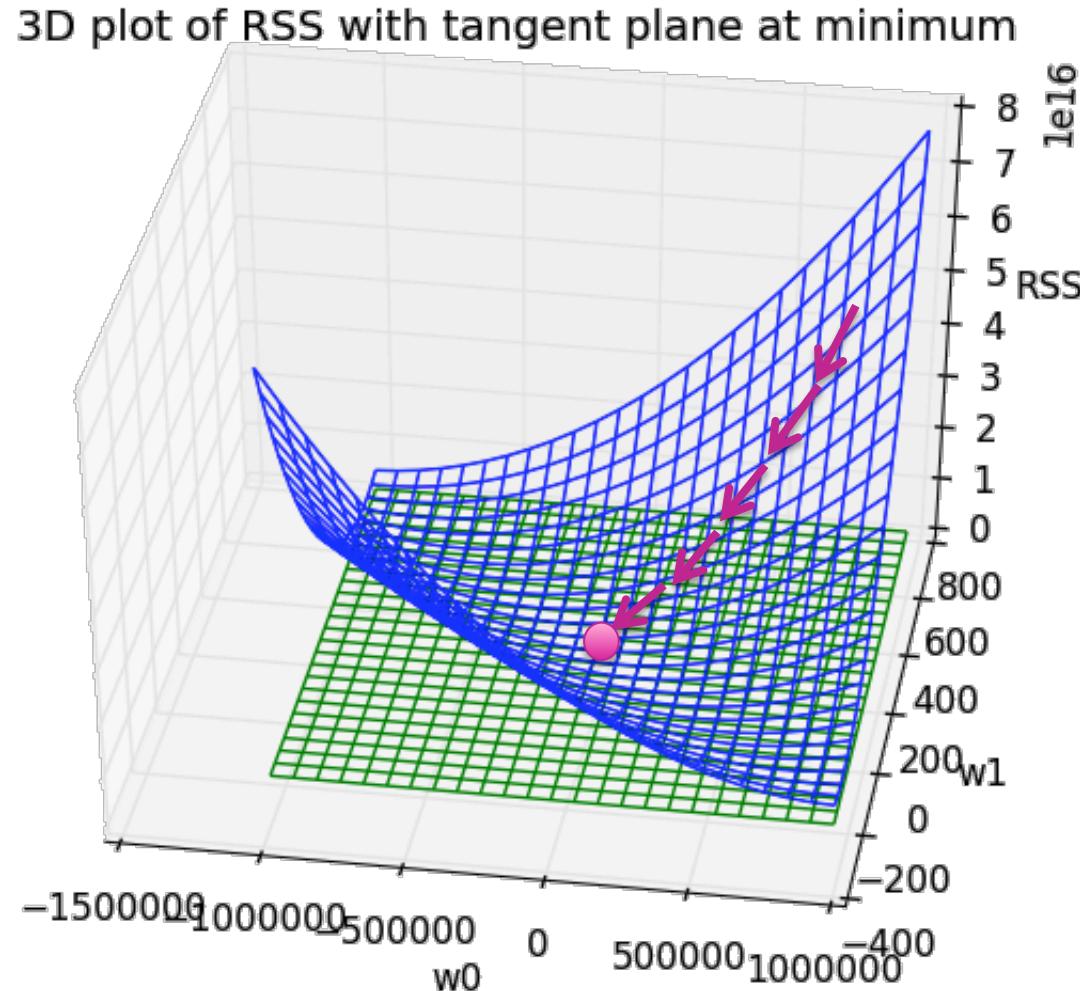


Minimize function  
over all possible  $w_0, w_1$

$$\min_{w_0, w_1} \sum_{i=1}^N (y_i - [w_0 + w_1 x_i])^2$$

RSS( $w_0, w_1$ ) is a function  
of 2 variables

# Gradient descent



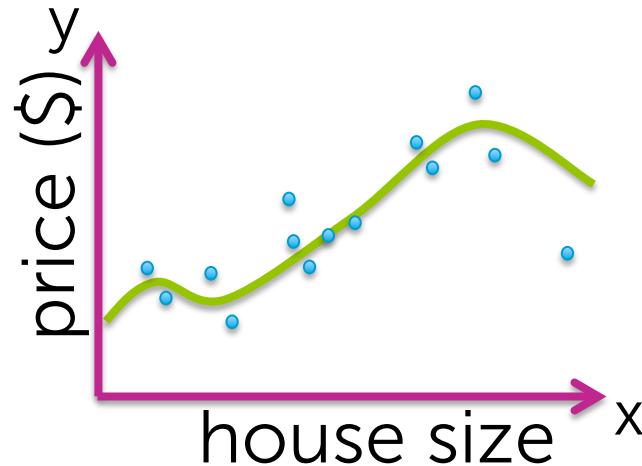
Algorithm:

**while** not converged

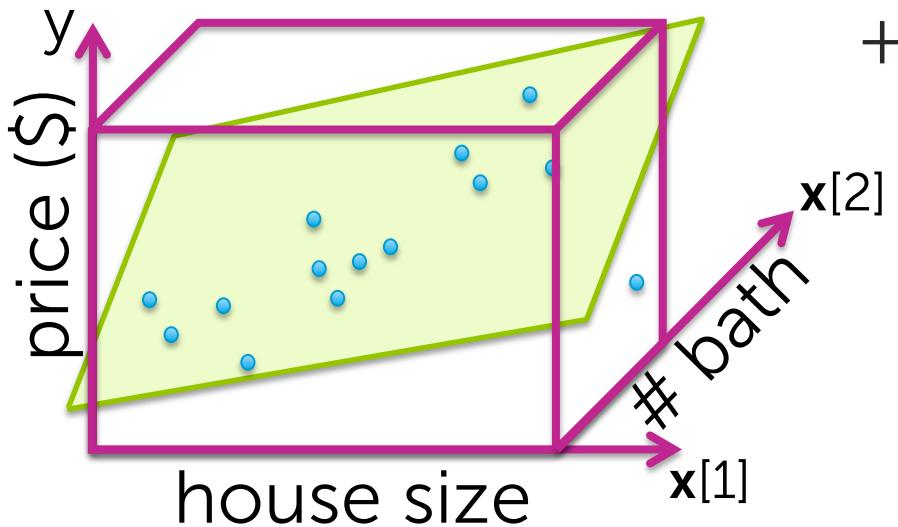
$$w^{(t+1)} \leftarrow w^{(t)} - \eta \nabla \text{RSS}(w^{(t)})$$

# Module 2: Multiple Regression

# Regression with multiple features



Fit more complex relationships than just a line



Incorporate more inputs + features thereof

- Square feet
- # bathrooms
- # bedrooms
- Lot size
- Year built
- ...

# Formally...

Model:

$$\begin{aligned}y_i &= \sum_{j=0}^D w_j h_j(\mathbf{x}_i) + \epsilon_i \\&= \sum_{j=0}^D w_j h_j(\mathbf{x}_i) + \epsilon_i\end{aligned}$$

*feature 1* =  $h_0(\mathbf{x})$  ... e.g., 1

*feature 2* =  $h_1(\mathbf{x})$  ... e.g.,  $\mathbf{x}[1]$  = sq. ft.

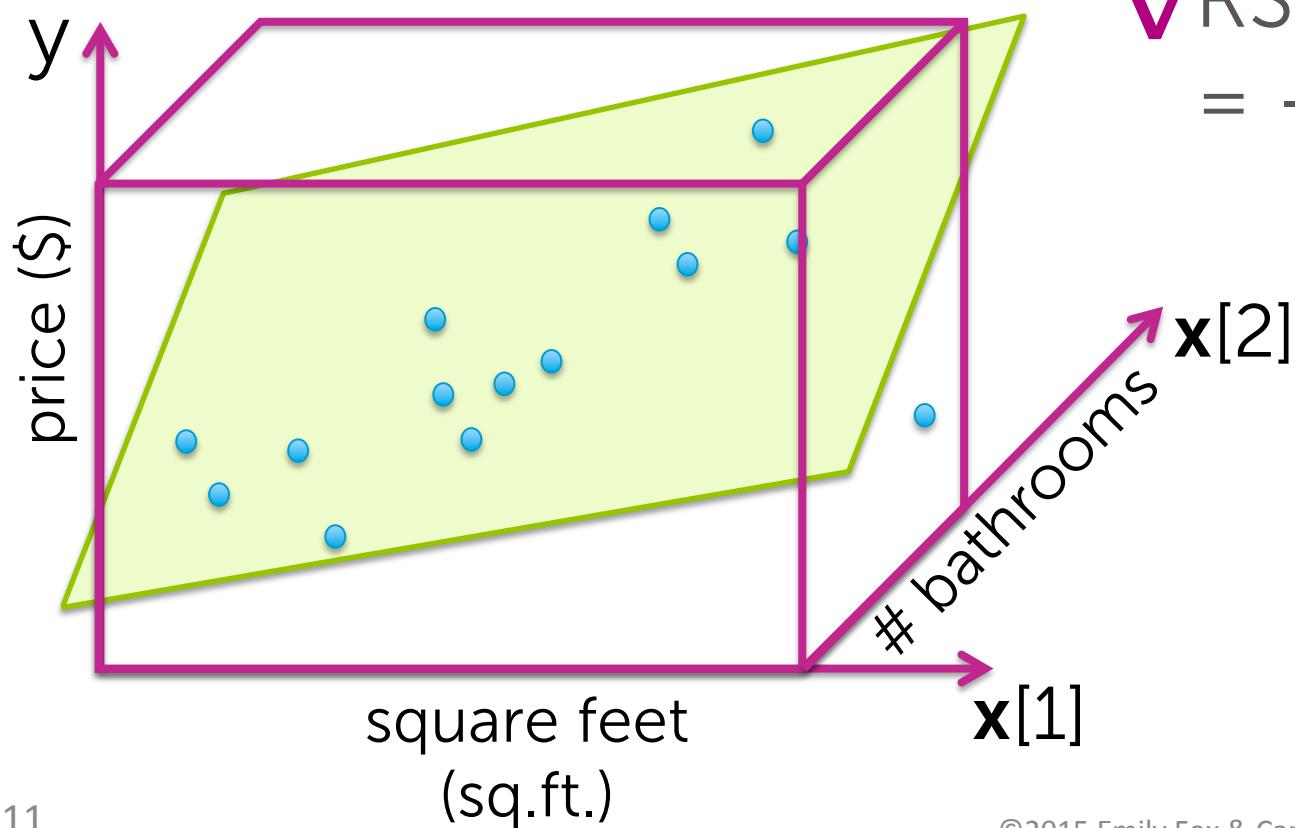
*feature 3* =  $h_2(\mathbf{x})$  ... e.g.,  $\mathbf{x}[2]$  = #bath  
or,  $\log(\mathbf{x}[7]) \mathbf{x}[2] = \log(\#bed) \times \#bath$

...

*feature D+1* =  $h_D(\mathbf{x})$  ... some other function of  $\mathbf{x}[1], \dots, \mathbf{x}[d]$

# RSS for multiple regression

$$\text{RSS}(\mathbf{w}) = \sum_{i=1}^N (y_i - h(x_i)^\top \mathbf{w})^2$$



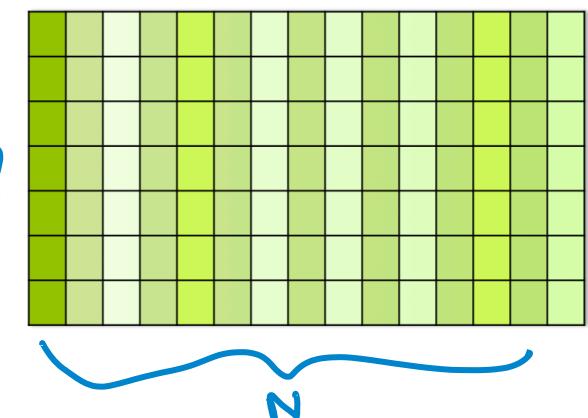
$$\nabla \text{RSS}(\mathbf{w}) = -2\mathbf{H}^\top(\mathbf{y} - \mathbf{H}\mathbf{w})$$

# Closed-form solution

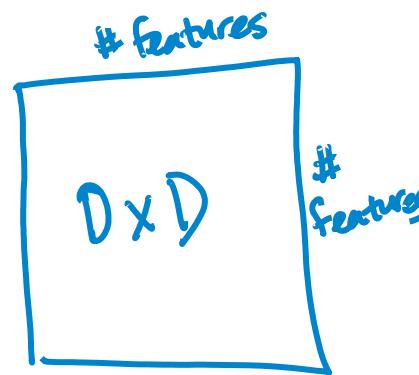
$$\hat{\mathbf{w}} = (\mathbf{H}^T \mathbf{H})^{-1} \mathbf{H}^T \mathbf{y}$$



# features = D



# obs = N



Invertible if:

In most cases

is  $N > D$

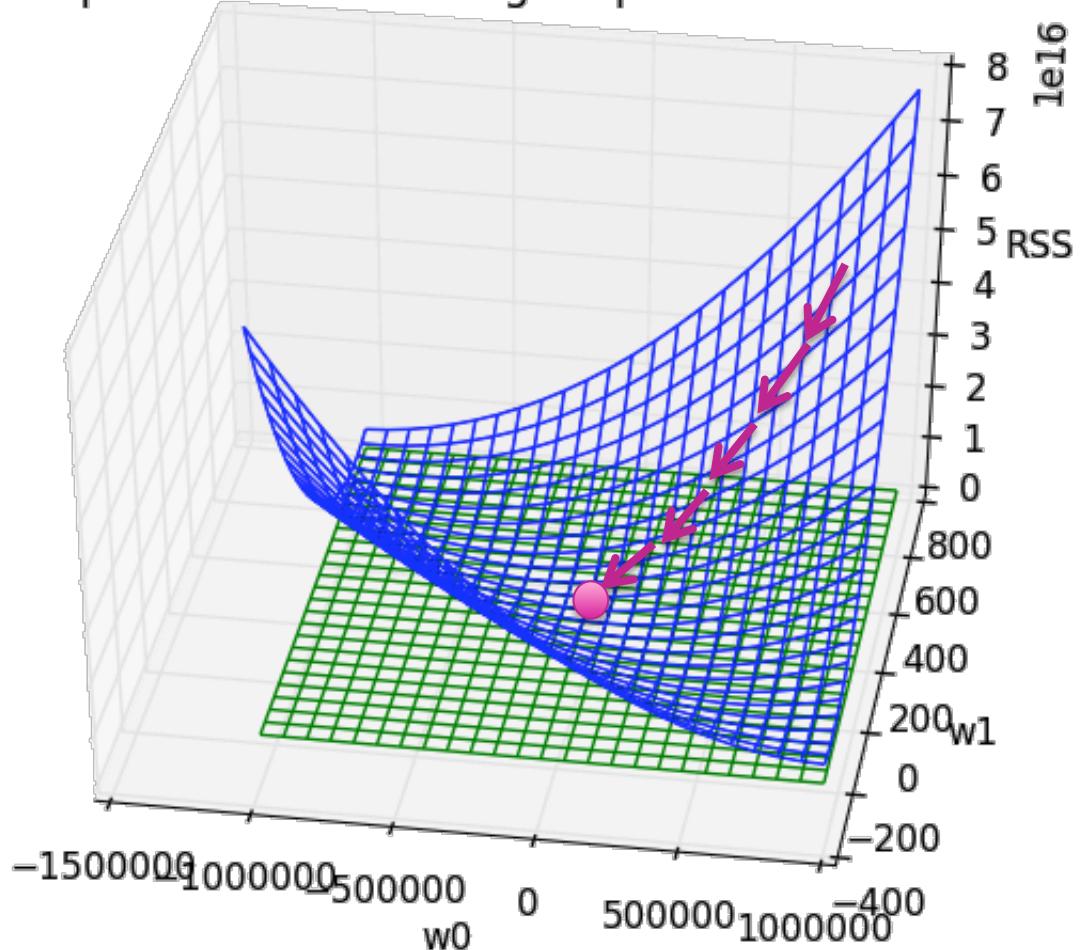
really,  
# of linearly  
ind. observations

Complexity of inverse:

$O(D^3)$

# Gradient descent for multiple regression

3D plot of RSS with tangent plane at minimum



init  $\mathbf{w}^{(1)} = 0$  (or randomly, or smartly),  $t=1$

**while**  $\|\nabla \text{RSS}(\mathbf{w}^{(t)})\| > \epsilon$

**for**  $j=0, \dots, D$

$\text{partial}[j] = -2 \sum_{i=1}^N h_j(\mathbf{x}_i)(y_i - \hat{y}_i(\mathbf{w}^{(t)}))$

$\mathbf{w}_j^{(t+1)} \leftarrow \mathbf{w}_j^{(t)} - \eta \text{ partial}[j]$

$t \leftarrow t + 1$

# Module 3: Assessing Performance

# Measuring loss

Loss function:

Cost of using  $\hat{w}$  at  $x$   
when  $y$  is true

$$L(y, f_{\hat{w}}(\mathbf{x}))$$

actual value  $\hat{f}(\mathbf{x}) = \text{predicted value } \hat{y}$

Examples:

(assuming loss for underpredicting = overpredicting)

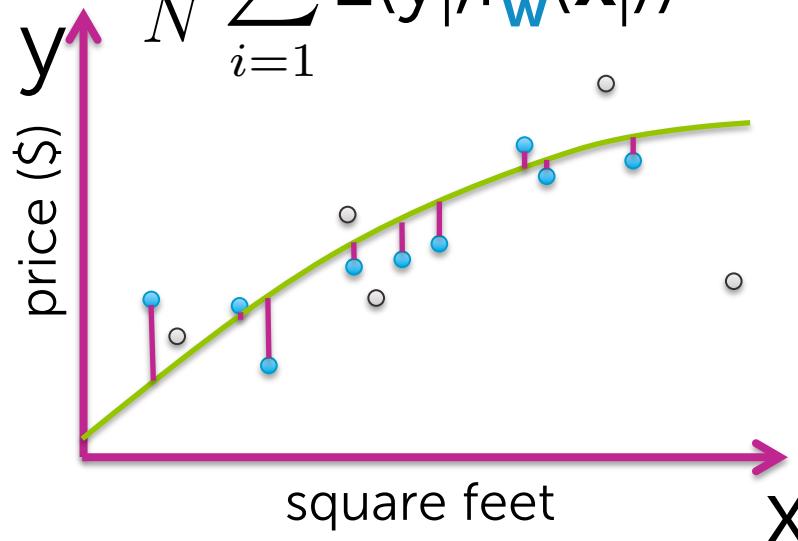
Absolute error:  $L(y, f_{\hat{w}}(\mathbf{x})) = |y - f_{\hat{w}}(\mathbf{x})|$

Squared error:  $L(y, f_{\hat{w}}(\mathbf{x})) = (y - f_{\hat{w}}(\mathbf{x}))^2$

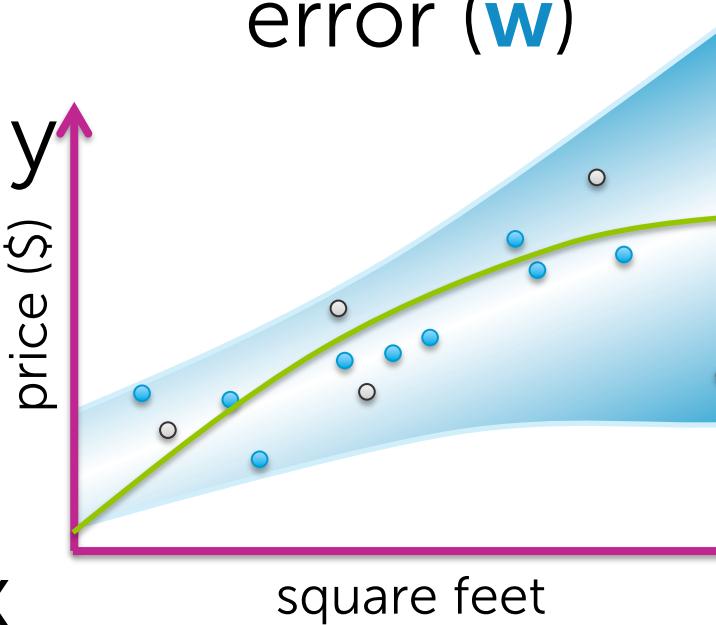
# 3 Measures of error

Training error ( $\hat{w}$ ) =

$$\frac{1}{N} \sum_{i=1}^N L(y_i, f_{\hat{w}}(\mathbf{x}_i))$$

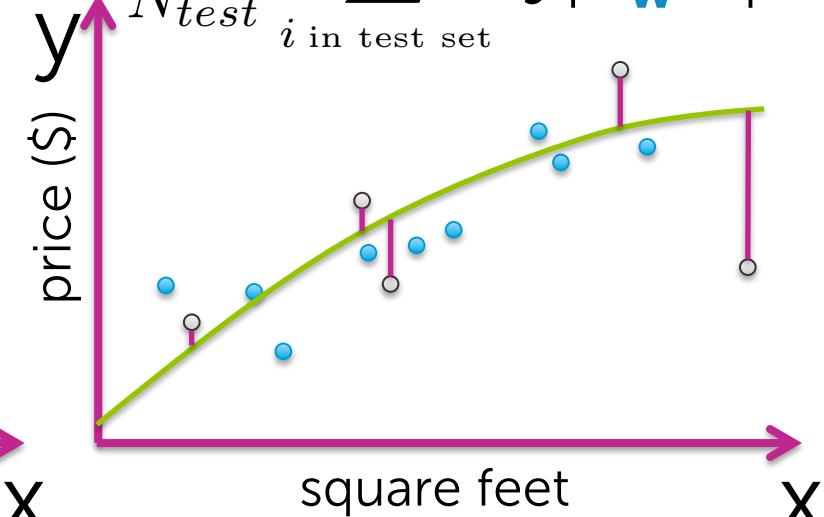


Generalization  
error ( $\hat{w}$ )

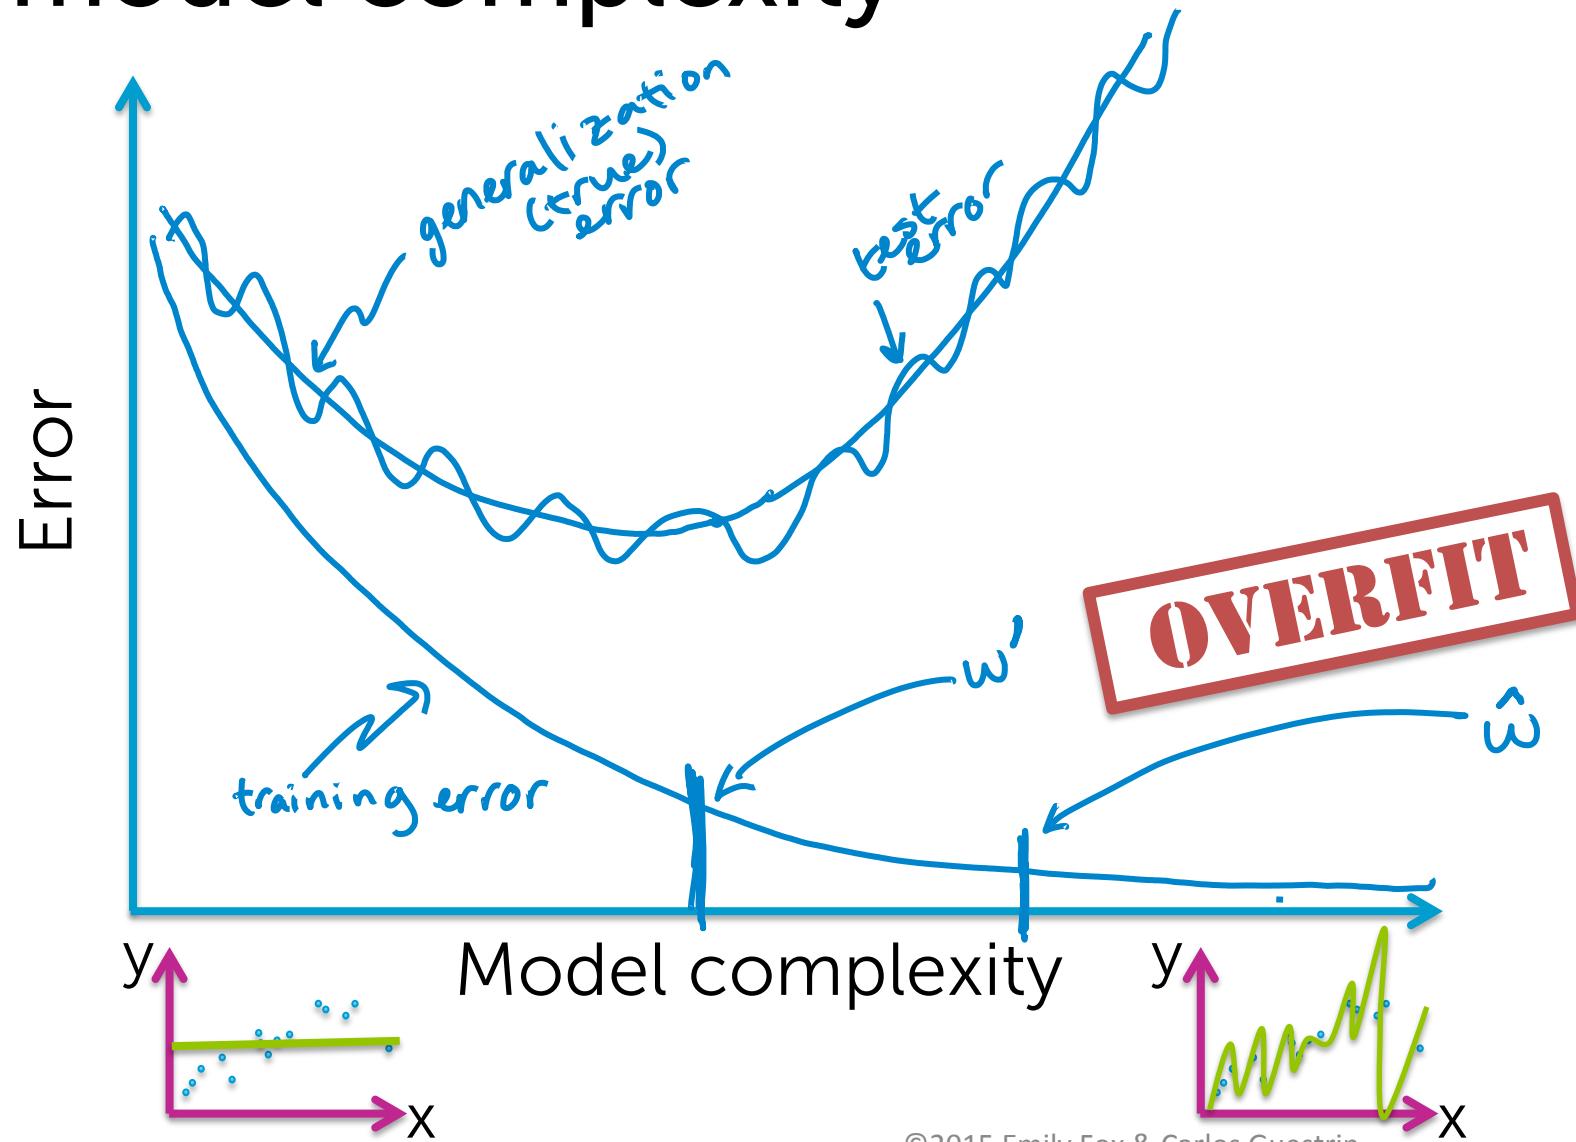


Test error ( $\hat{w}$ ) =

$$\frac{1}{N_{test}} \sum_{i \text{ in test set}} L(y_i, f_{\hat{w}}(\mathbf{x}_i))$$



# Training, true, & test error vs. model complexity

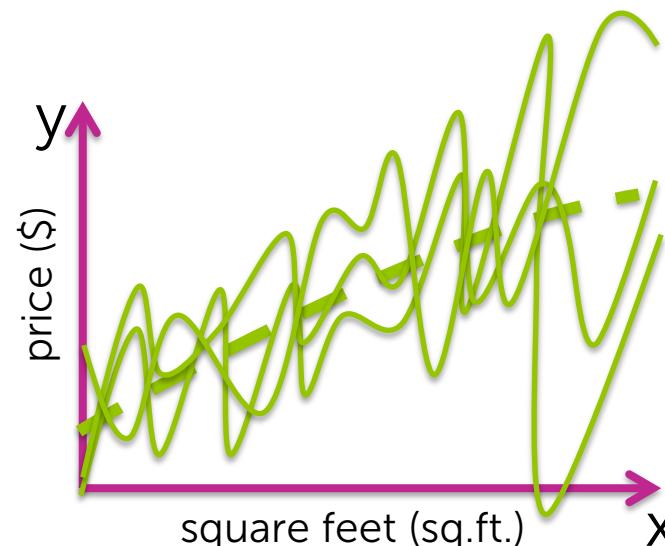
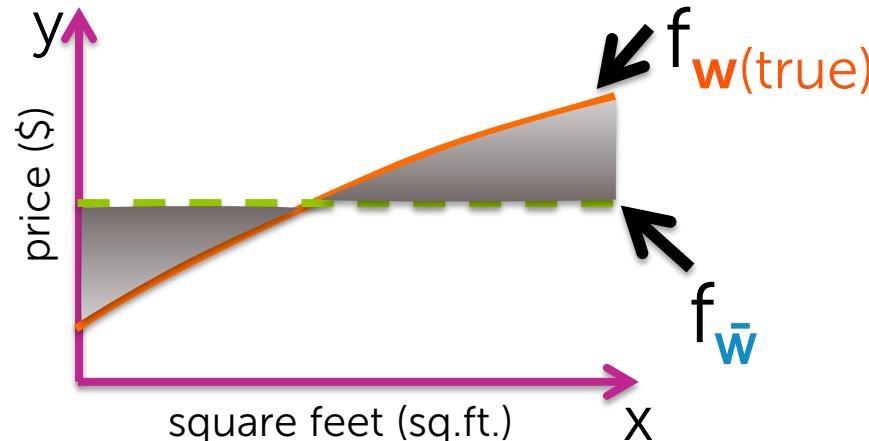


# 3 Sources of prediction error

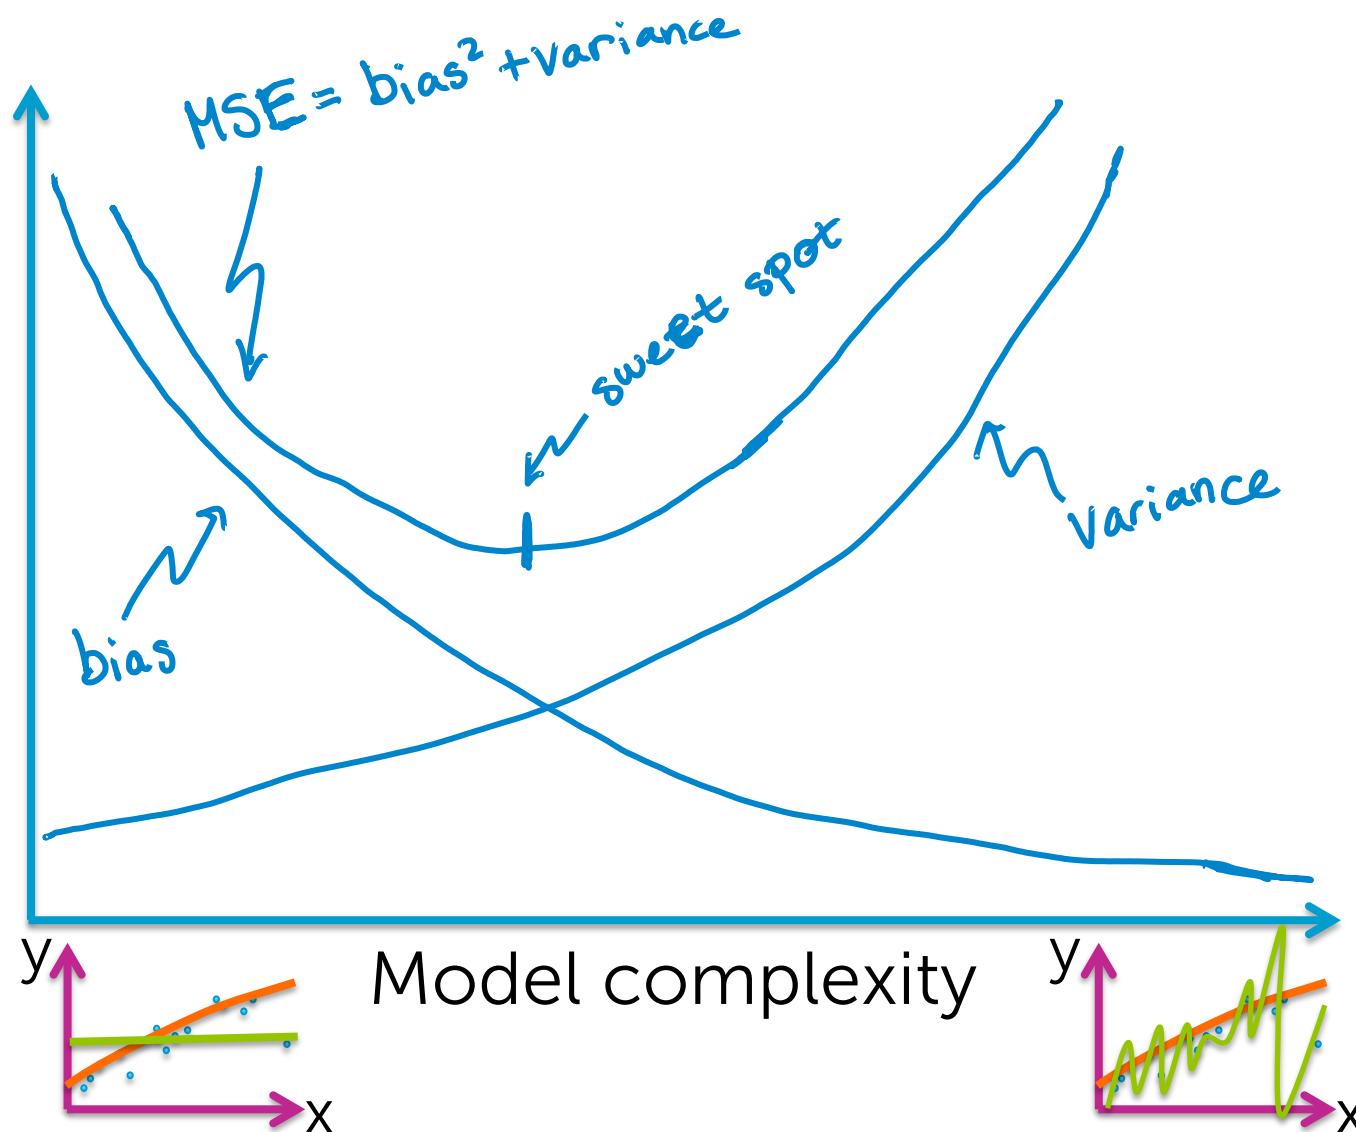
1. Noise

2. Bias

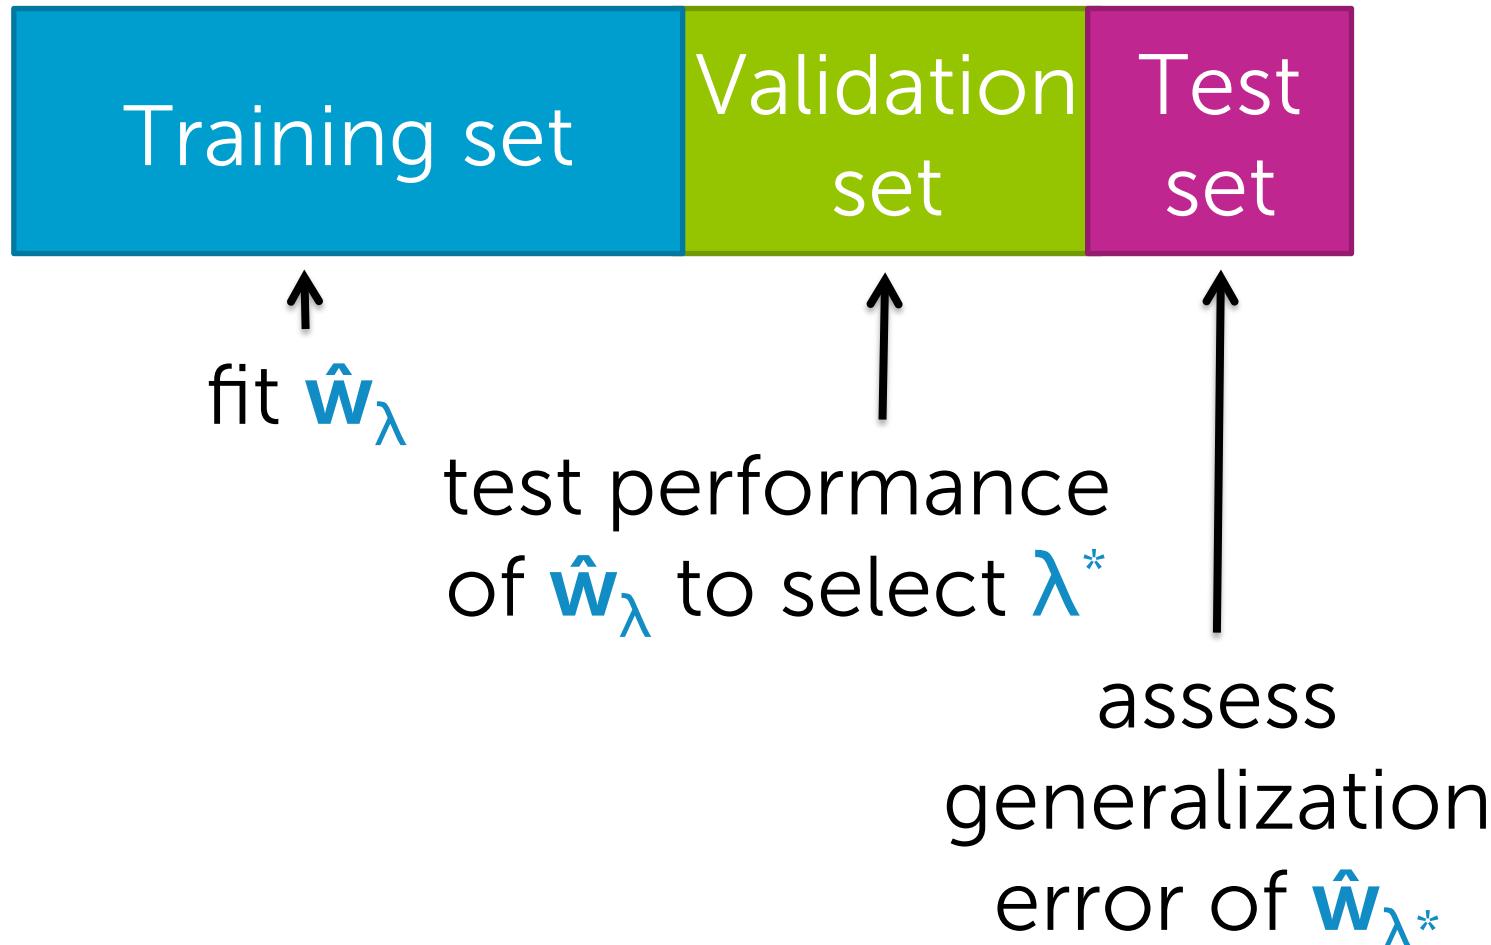
3. Variance



# Bias-variance tradeoff



# Model selection & assessment



# Module 4: Ridge Regression

# Balancing fit and model complexity



Ridge total cost =  
measure of fit + measure of magnitude  
of coefficients

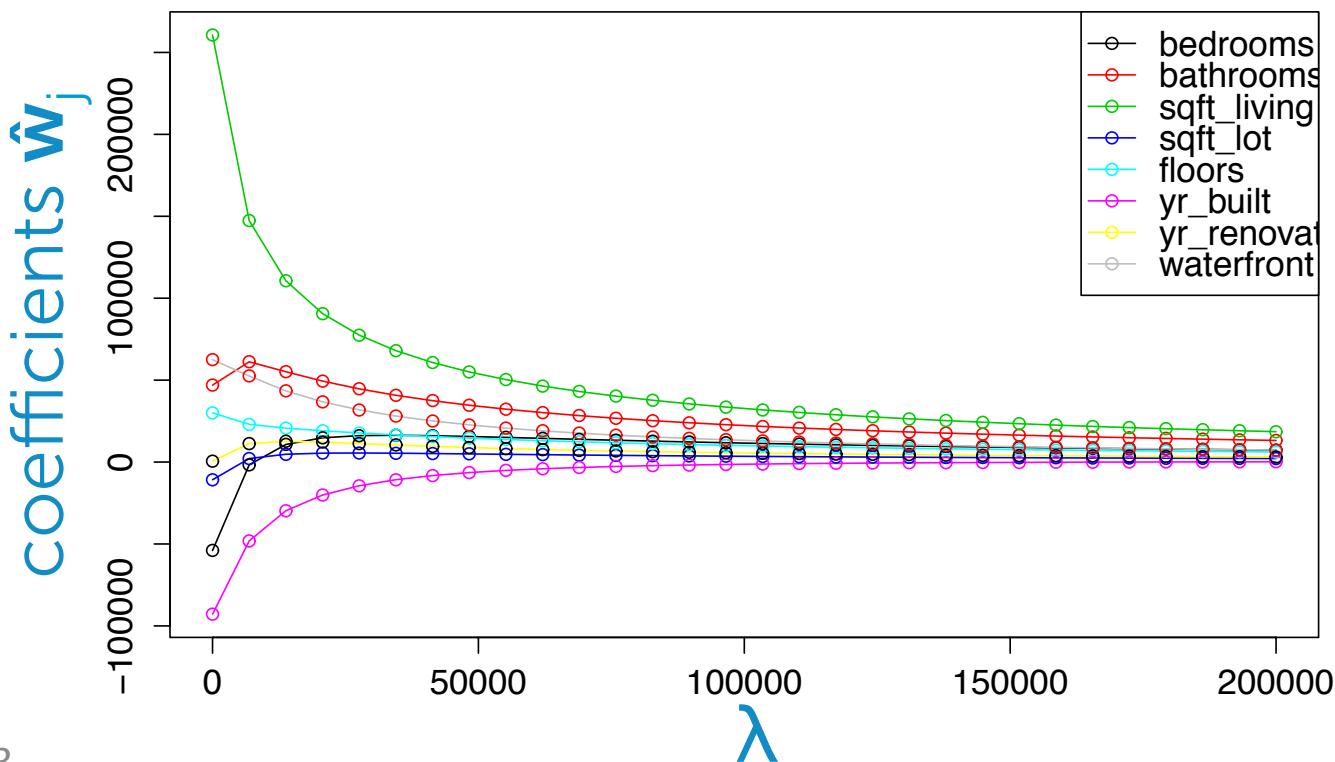
bias-variance tradeoff

# Ridge objective function ( $L_2$ regularized regression)

$\hat{\mathbf{w}}$  selected to minimize

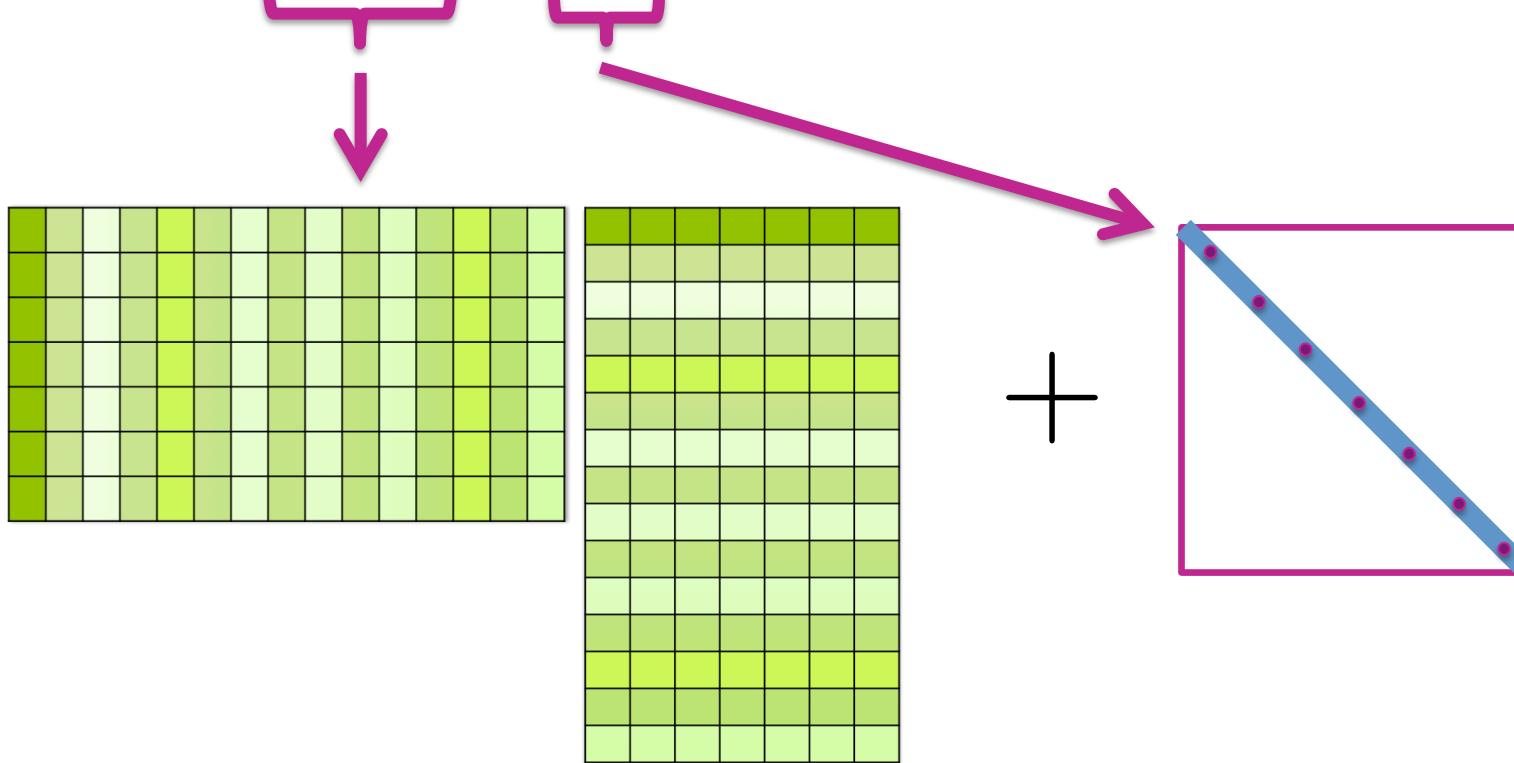
$$\text{RSS}(\mathbf{w}) + \lambda \|\mathbf{w}\|_2^2$$

tuning parameter =  
balance of fit and magnitude



# Ridge closed-form solution

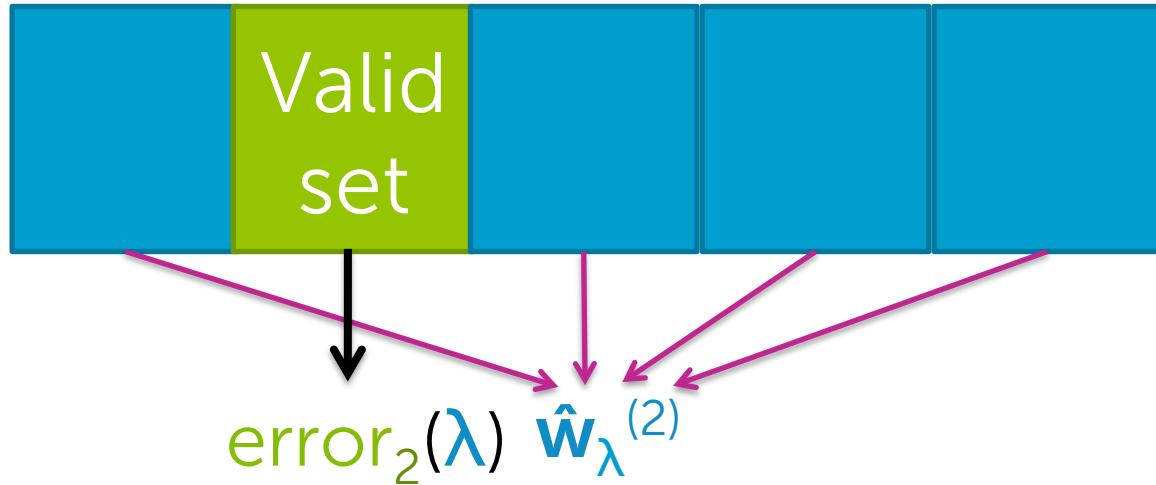
$$\hat{\mathbf{w}} = (\mathbf{H}^T \mathbf{H} + \lambda \mathbf{I})^{-1} \mathbf{H}^T \mathbf{y}$$



Invertible if:  
Always if  $\lambda > 0$ ,  
even if  $N < D$

Complexity of  
inverse:  
 $O(D^3)$ ...  
big for large  $D$ !

# K-fold cross validation



For  $k=1, \dots, K$

1. Estimate  $\hat{w}_\lambda^{(k)}$  on the training blocks
2. Compute error on validation block:  $\text{error}_k(\lambda)$

Compute average error:  $\text{CV}(\lambda) = \frac{1}{K} \sum_{k=1}^K \text{error}_k(\lambda)$

# Module 5: Lasso Regression

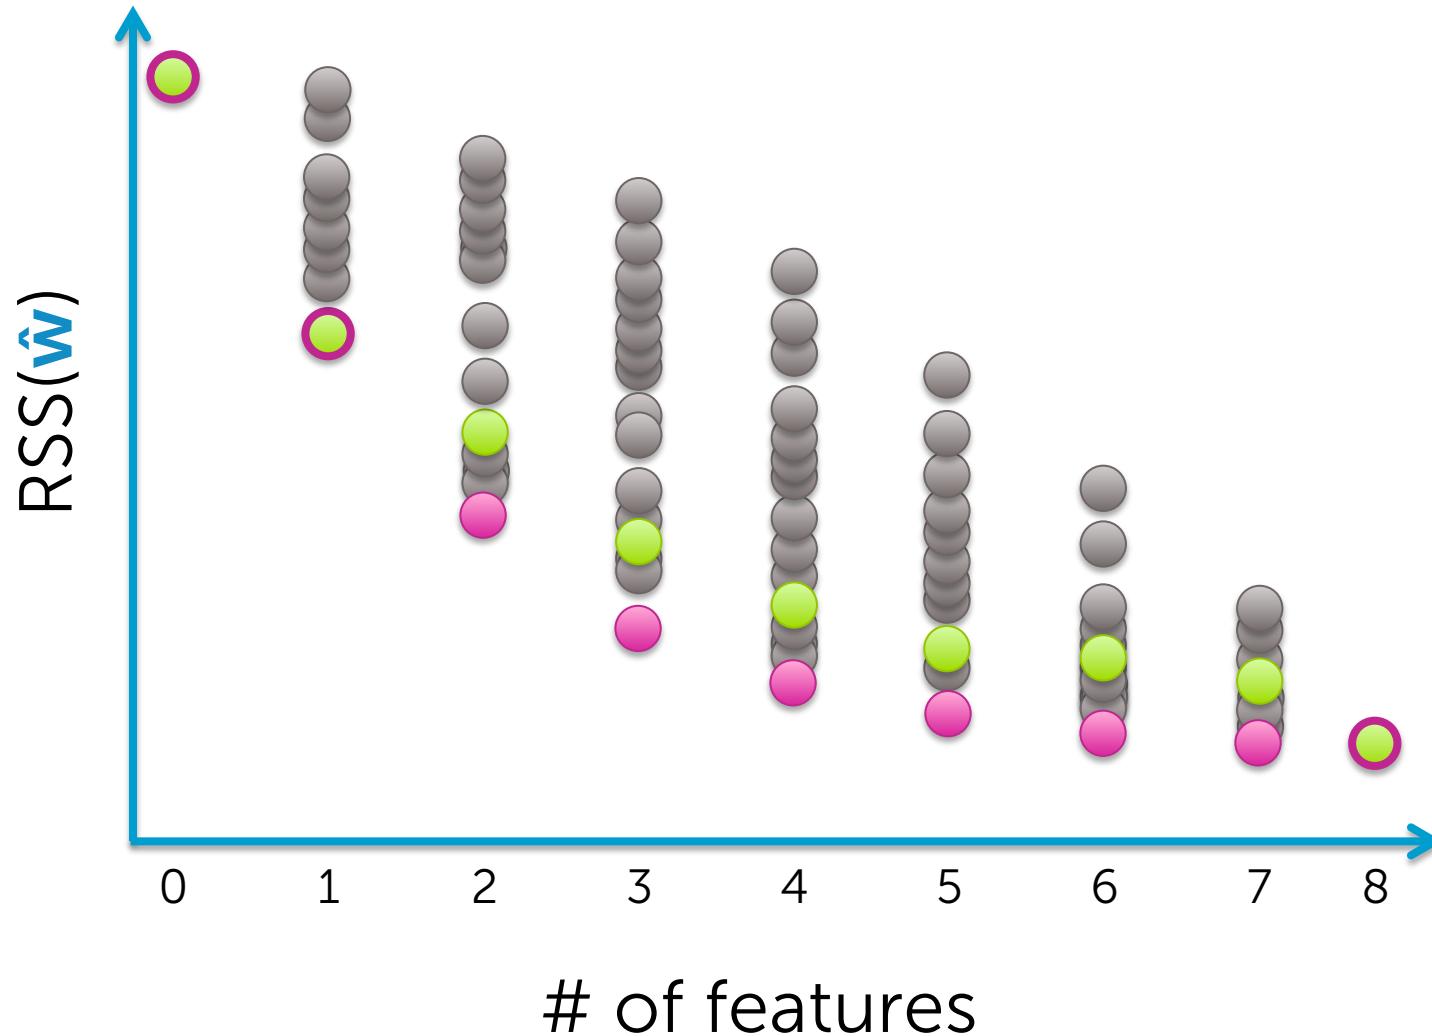
# Performing feature selection



Useful for **efficiency**  
of predictions and  
**interpretability**

Lot size	Dishwasher
Single Family	Garbage disposal
Year built	Microwave
Last sold price	Range / Oven
Last sale price/sqft	Refrigerator
Finished sqft	Washer
Unfinished sqft	Dryer
Finished basement sqft	Laundry location
# floors	Heating type
Flooring types	Jetted Tub
Parking type	Deck
Parking amount	Fenced Yard
Cooling	Lawn
Heating	Garden
Exterior materials	Sprinkler System
Roof type	:
Structure style	:

# All subsets vs. greedy



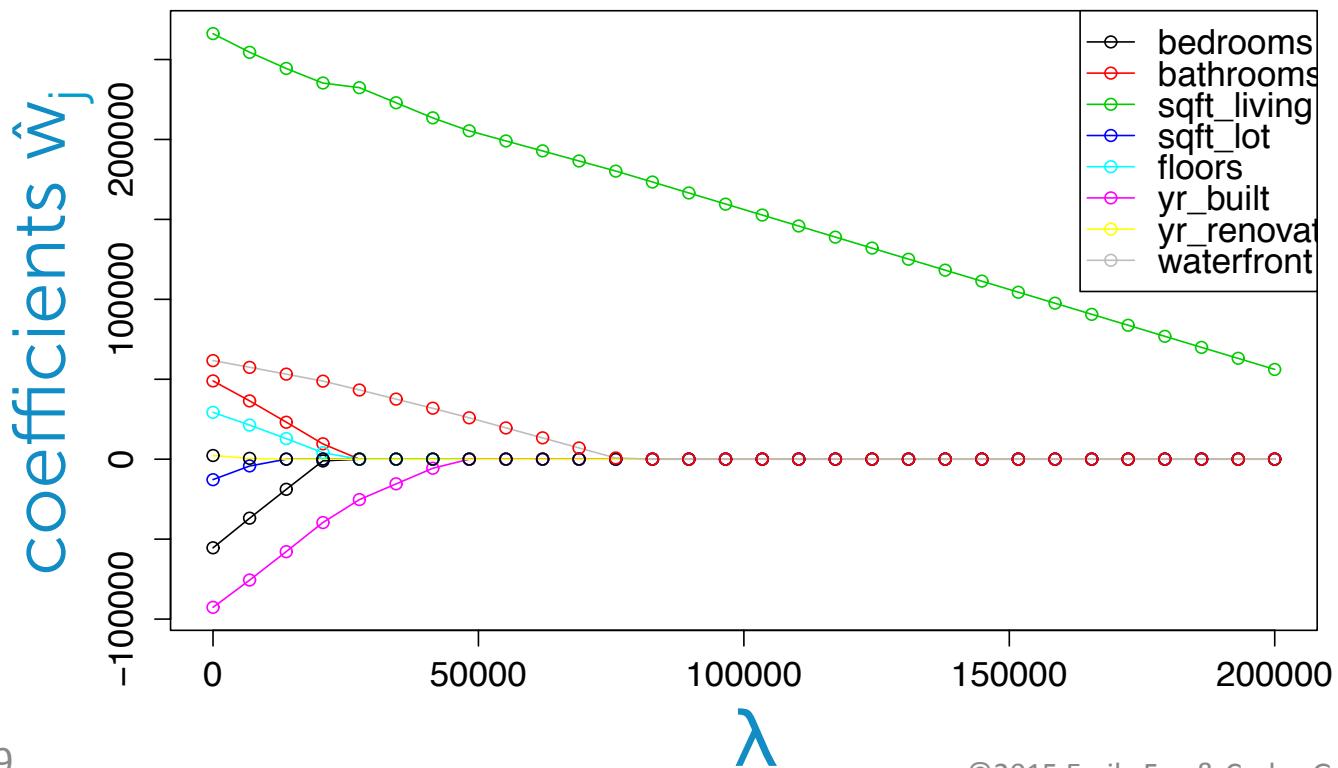
- # bedrooms
- # bathrooms
- sq.ft. living
- sq.ft. lot
- floors
- year built
- year renovated
- waterfront

# Lasso objective function ( $L_1$ regularized regression)

$\hat{\mathbf{w}}$  selected to minimize

$$\text{RSS}(\mathbf{w}) + \lambda \|\mathbf{w}\|_1$$

tuning parameter =  
balance of fit and sparsity

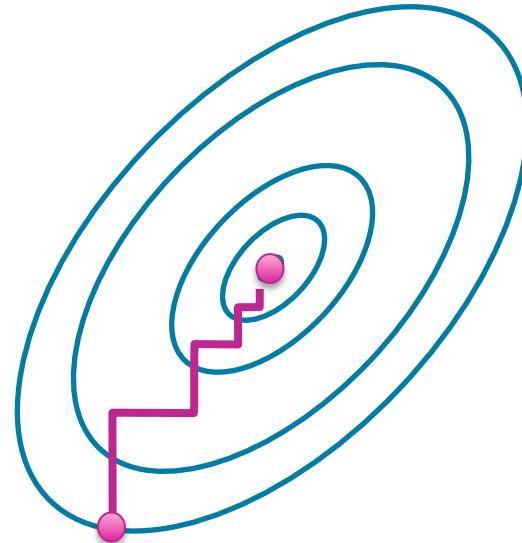


# Coordinate descent for lasso

Precompute:  $z_j = \sum_{i=1}^N h_j(\mathbf{x}_i)^2$

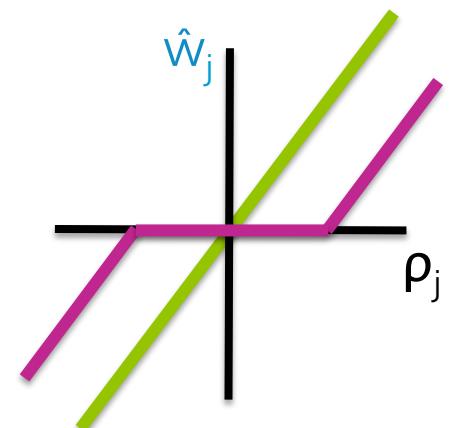
Initialize  $\hat{\mathbf{w}} = 0$  (or smartly...)  
while not converged

for  $j=0,1,\dots,D$



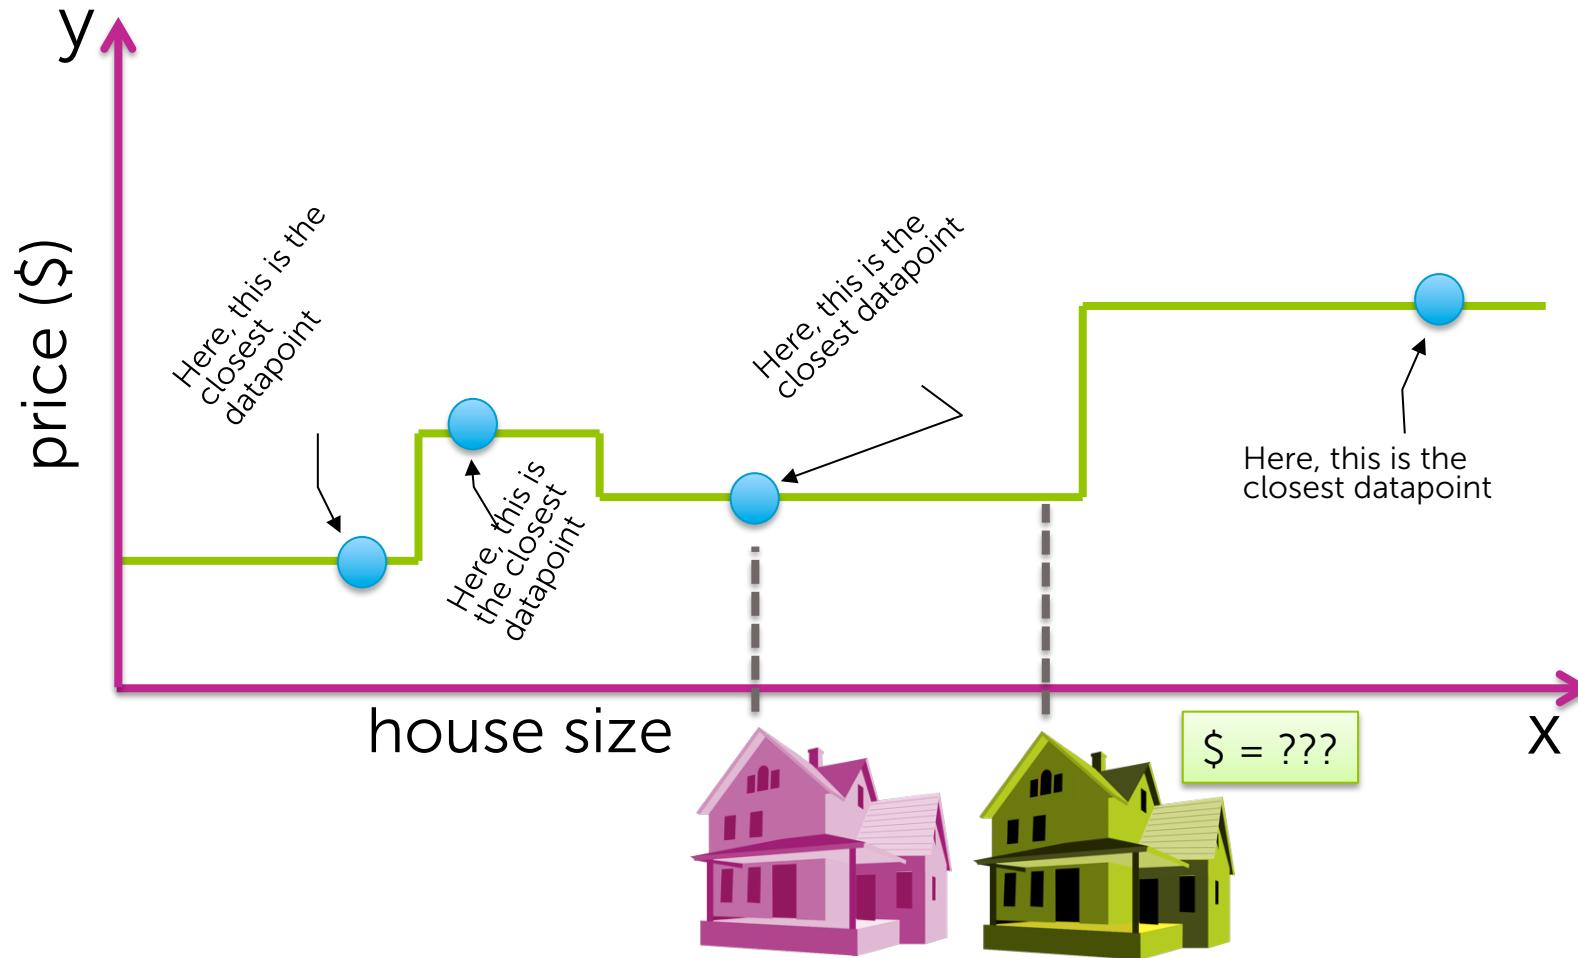
compute:  $\rho_j = \sum_{i=1}^N h_j(\mathbf{x}_i)(y_i - \hat{y}_i(\hat{\mathbf{w}}_{-j}))$

set:  $\hat{w}_j = \begin{cases} (\rho_j + \lambda/2)/z_j & \text{if } \rho_j < -\lambda/2 \\ 0 & \text{if } \rho_j \text{ in } [-\lambda/2, \lambda/2] \\ (\rho_j - \lambda/2)/z_j & \text{if } \rho_j > \lambda/2 \end{cases}$



# Module 6: Nearest Neighbor & Kernel Regression

# 1-Nearest neighbor regression



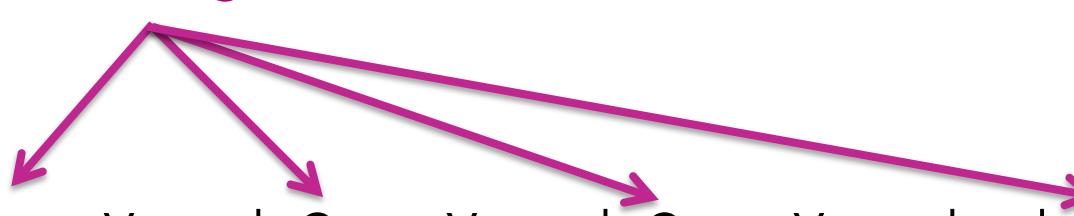
# Weighted k-NN

Weigh more similar houses more than those less similar in list of k-NN

Predict:

$$\hat{y}_q = \frac{c_{qNN1}y_{NN1} + c_{qNN2}y_{NN2} + c_{qNN3}y_{NN3} + \dots + c_{qNNk}y_{NNk}}{\sum_{j=1}^k c_{qNNj}}$$

weights on NN



# Kernel regression

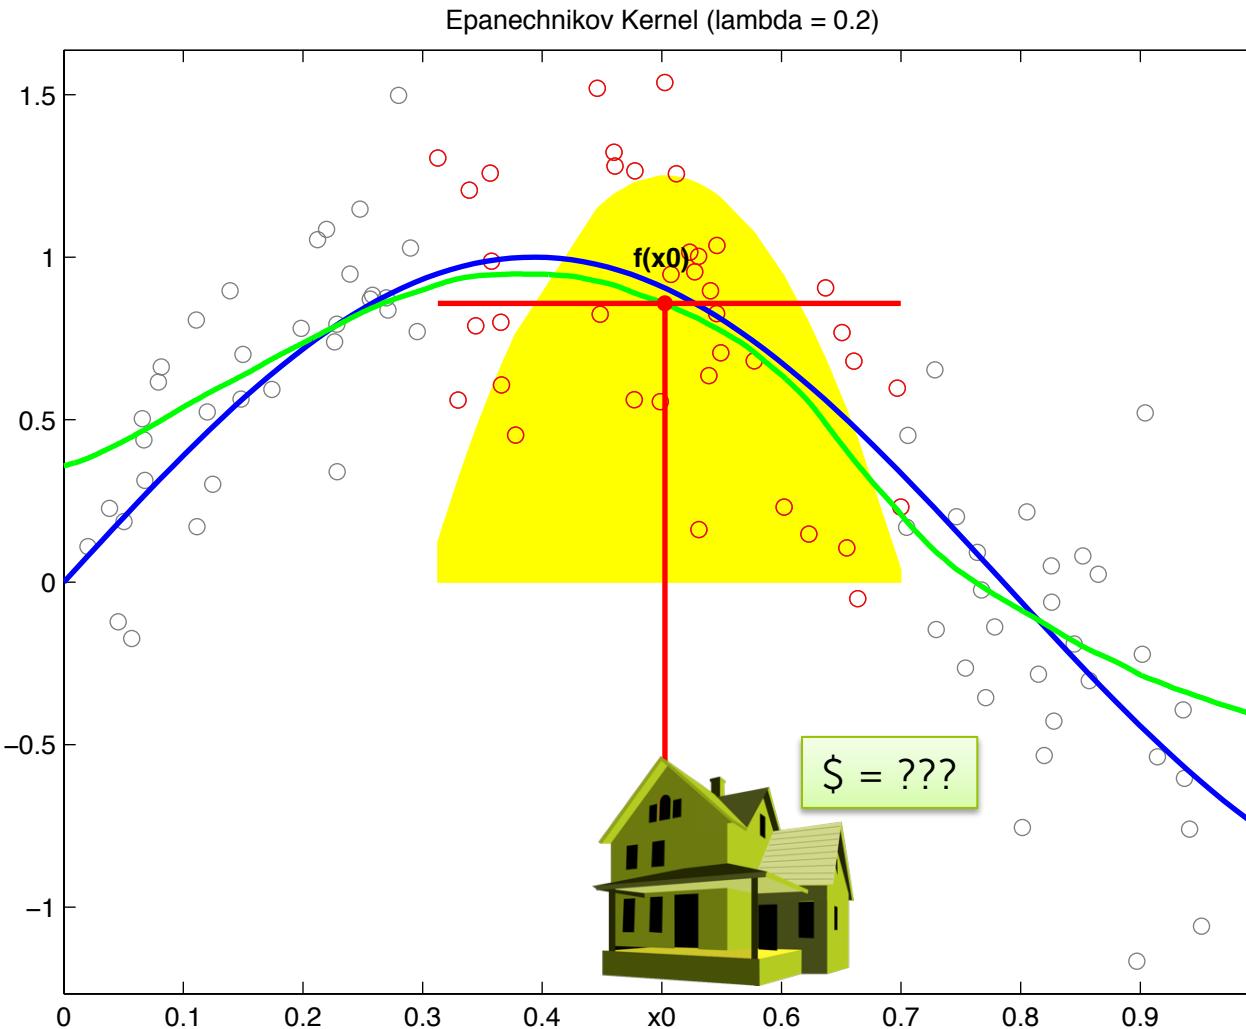
Instead of just weighting NN, weight *all* points

Predict:

weight on each datapoint

$$\hat{y}_q = \frac{\sum_{i=1}^N c_{qi} y_i}{\sum_{i=1}^N c_{qi}} = \frac{\sum_{i=1}^N \text{Kernel}_{\lambda}(\text{distance}(\mathbf{x}_i, \mathbf{x}_q)) * y_i}{\sum_{i=1}^N \text{Kernel}_{\lambda}(\text{distance}(\mathbf{x}_i, \mathbf{x}_q))}$$

# Visualizing kernel regression



# Summary of what we learned

## Models

- Linear regression
- Regularization: Ridge (L2), Lasso (L1)
- Nearest neighbor and kernel regression

## Algorithms

- Gradient descent
- Coordinate descent

## Concepts

- Loss functions, bias-variance tradeoff, cross-validation, sparsity, overfitting, model selection, feature selection

# What we didn't cover

# Other important regression topics

- Multivariate outputs  $\mathbf{y}$ 
  - ...when correlated
- Maximum likelihood estimation
  - Equivalent to least squares when errors are “normal”/Gaussian
- Statistical inferences
- “Generalized linear models”
  - Models for non-Gaussian error
  - E.g., outputs are
    - (i) constrained to be positive or bounded
    - (ii) discrete (“yes”/“no”)
- Regression trees
- ...

# What's ahead in this specialization

# 3. Classification

## *Case study: Analyzing sentiment*

Models

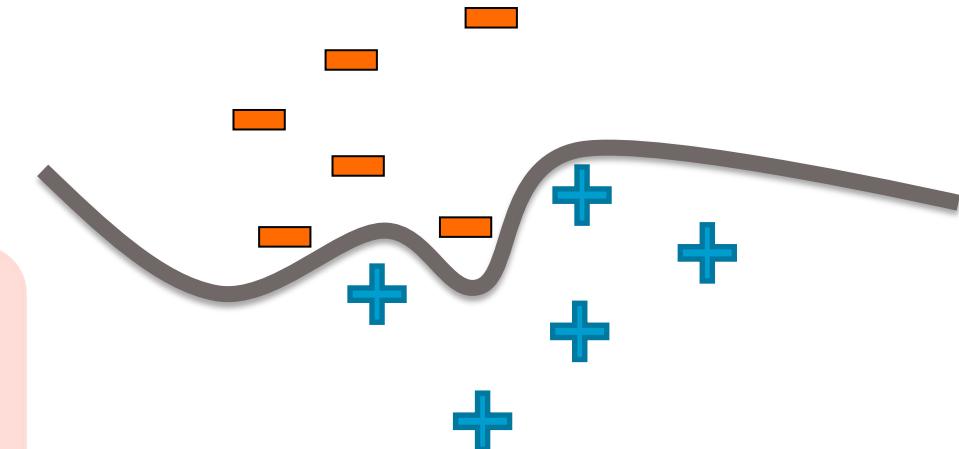
- Linear classifiers
- Decision trees
- Boosted trees and random forests

Algorithms

- Stochastic gradient descent
- Boosting

Concepts

- Decision boundaries, MLE, ensemble methods, online learning



# 4. Clustering & Retrieval

## *Case study: Finding documents*

Models

- Nearest neighbors
- Clustering, mixtures of Gaussians
- Latent Dirichlet allocation (LDA)

Algorithms

- KD-trees
- K-means
- Expectation-maximization (EM)

Concepts

- Distance metrics, approximation algorithms, sampling algorithms, scaling up with map-reduce



SPORTS



WORLD NEWS



ENTERTAINMENT



SCIENCE

# 5. Recommender Systems & Dimensionality Reduction

*Case study: Recommending Products*

Models

- Collaborative filtering
- Matrix factorization
- PCA

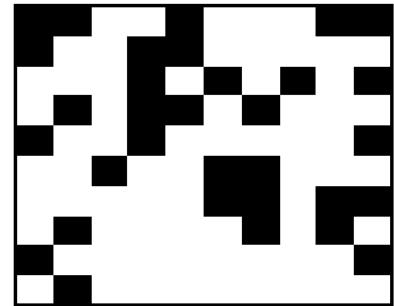
Algorithms

- Coordinate descent
- Eigen decomposition
- SVD

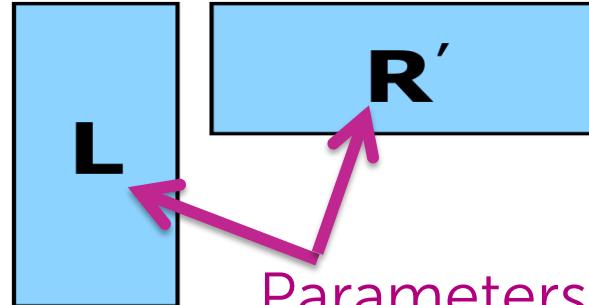
Concepts

- Matrix completion, eigenvalues, cold-start problem, diversity, scaling up

Rating =

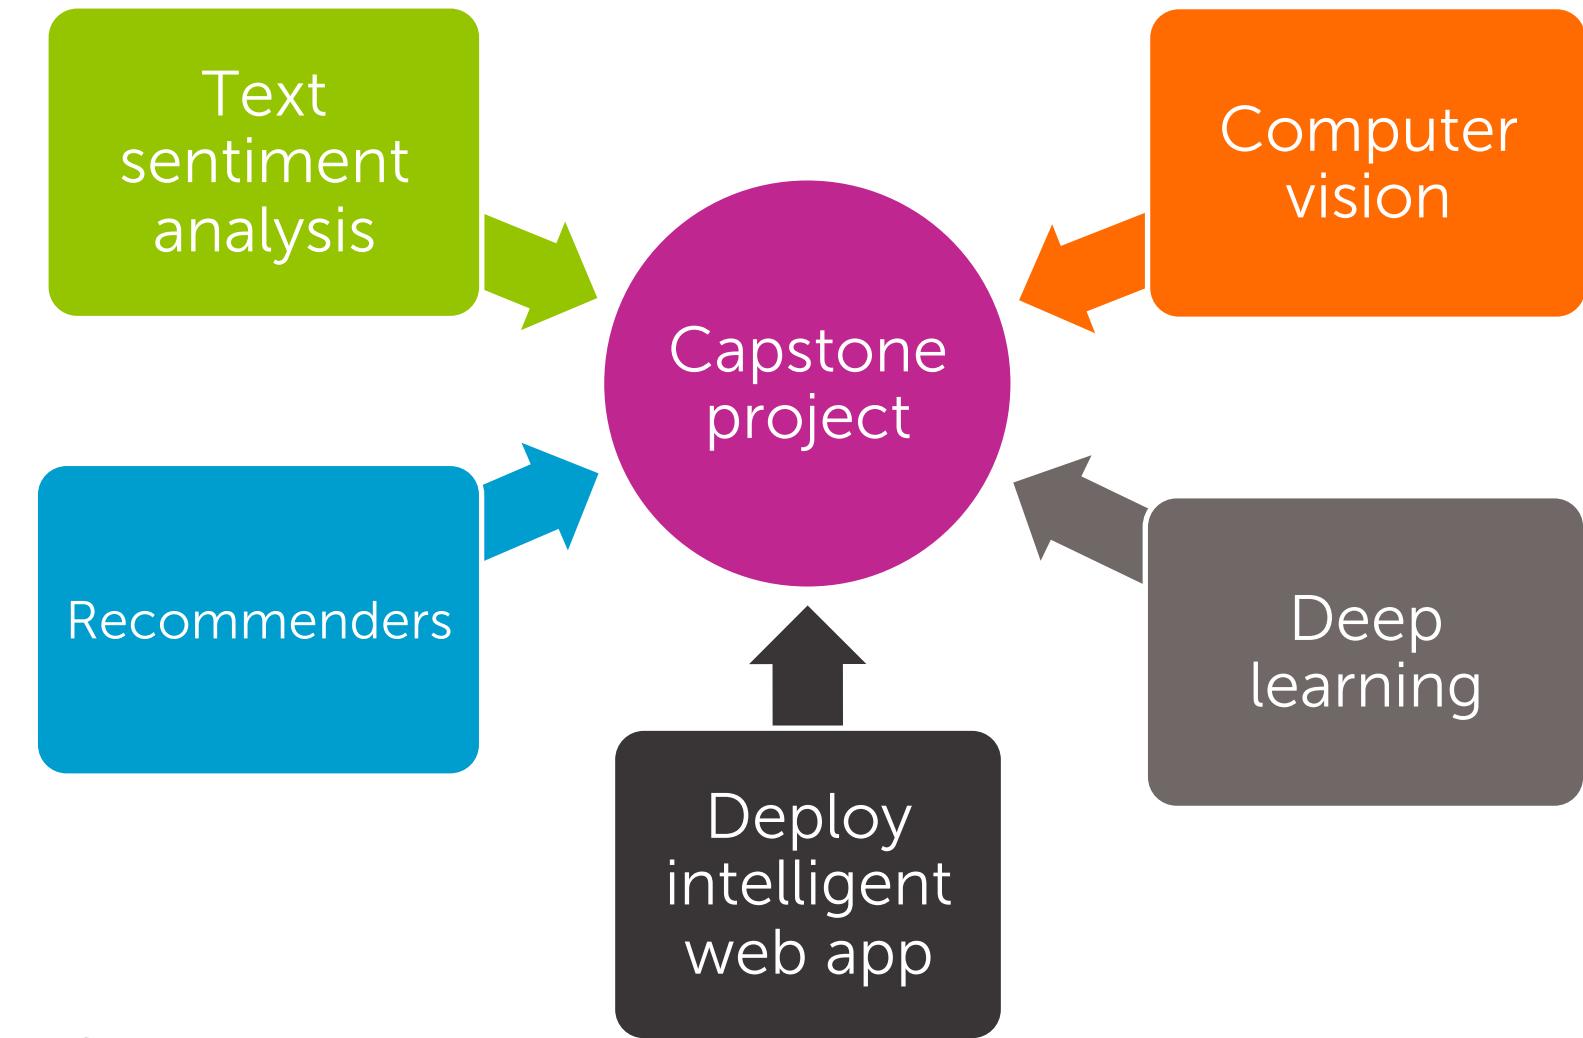


$\approx$



Parameters  
of model

# 6. Capstone: *Build and deploy an intelligent application with deep learning*





Thank you...